



`scripts/deploy`

REPOS

[**github.com/GhostX94/prometheus**](https://github.com/GhostX94/prometheus)

[**github.com/grobie/prometheus-on-kubernetes**](https://github.com/grobie/prometheus-on-kubernetes)

Monitoring Kubernetes with Prometheus

Isaac Rubio T, Jul 2018

Bootkube or Minikube

`scripts/deploy`

Prometheus

Monitoring system and time series database

Motivation

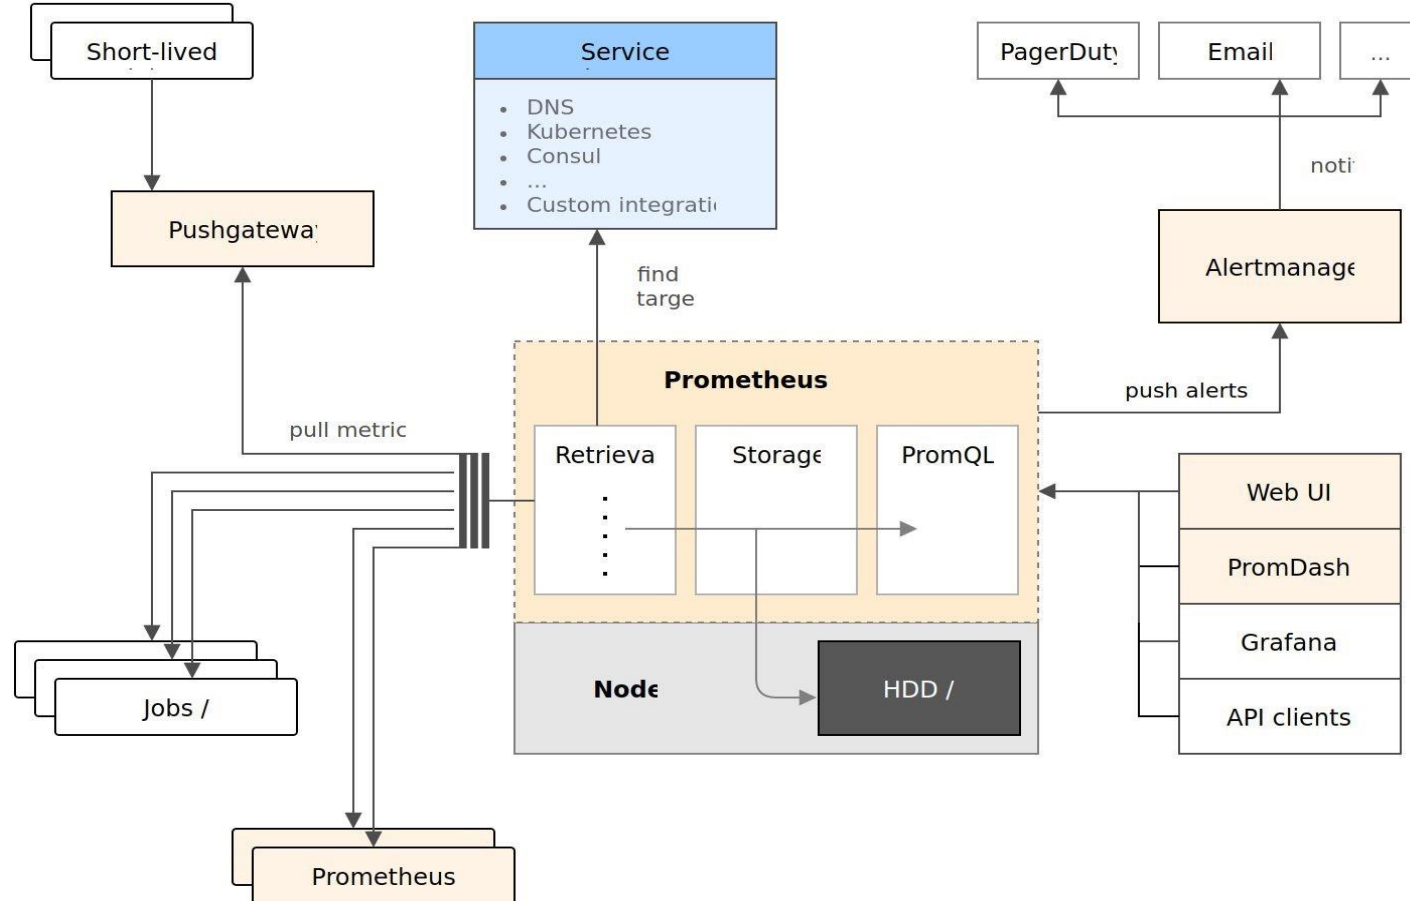
- Microservice architecture
 - Many more services than in traditional host-based monitoring
 - Short lifecycles
 - Heterogeneous workloads
- Insight
 - Detailed (instance/endpoint/version/... drilldown) and aggregated

(across a service) of everything (hardware to service)

- Trends (act before something becomes a problem)
- Alerting
 - Symptom vs. Cause
 - Grouping, flexible silencing

scripts/deploy

Overview



Examples

```
# HELP etcd_store_writes_total Total number of writes seen
...
# TYPE etcd_store_writes_total counter
etcd_store_writes_total{action="compareAndDelete"} 2
etcd_store_writes_total{action="compareAndSwap"} 4016
etcd_store_writes_total{action="create"} 218
etcd_store_writes_total{action="set"} 5
```

scripts/deploy

```
count by(job)(up == 0) / count by(job)(up)
rate(etcd_store_writes_total{action="set"}[1m]))
sum
without(action)(rate(etcd_store_writes_total[1m]))
```

Configuration

```
# prometheus.yaml
prometheus.io/docs/operating/configuration/ global:
    # Settings applying to all jobs

scrape_configs:
```



```
# Define different scrape jobs
```

```
Rules_files:
```

```
# Load files specifying rules to pre-calculate  
expressions # as well as alerts.
```

Configuration

```
scrape_configs: - job_name:  
etcd    static_configs:    -  
targets: ["172.17.4.51:2379"]
```

scripts/deploy

- job_name: kube-components
kubernetes_sd_configs:
 - role: endpointsrelabel_configs:
- # Custom filtering and label mapping

Configuration

```
# continued relabel_configs: - action: keep  
source_labels: [__meta_kubernetes_service_name]
```

```
regex: "kube-(.*)-prometheus-discovery" - action:
keep    source_labels:
[__meta_kubernetes_endpoint_port_name]    regex:
"prometheus" - action: replace    source_labels:
[__meta_kubernetes_service_name]    target_label:
job    regex: "(kube-.*)-prometheus-discovery"
```

Kubernetes

Container orchestration system **Domain objects**

scripts/deploy

- Pod
 - Group of one or more containers, share context and namespaces
 - Co-located and co-scheduled (allows for side-cars)
- Service
 - Logical set of Pods, stable access points
- Deployment
 - Declaration of the desired state (what to run, how to get there)
- Daemon / Pet / Replica sets
 - Definition of groups of pods (each node / stateful / stateless)
- ConfigMap
 - Configuration data / files (can be mounted in containers)

Workshop

Monitoring Kubernetes with Prometheus

git checkout 1.setup

kubectl get nodes

Running Prometheus in Kubernetes

Installation and configuration

scripts/deploy

Running Prometheus

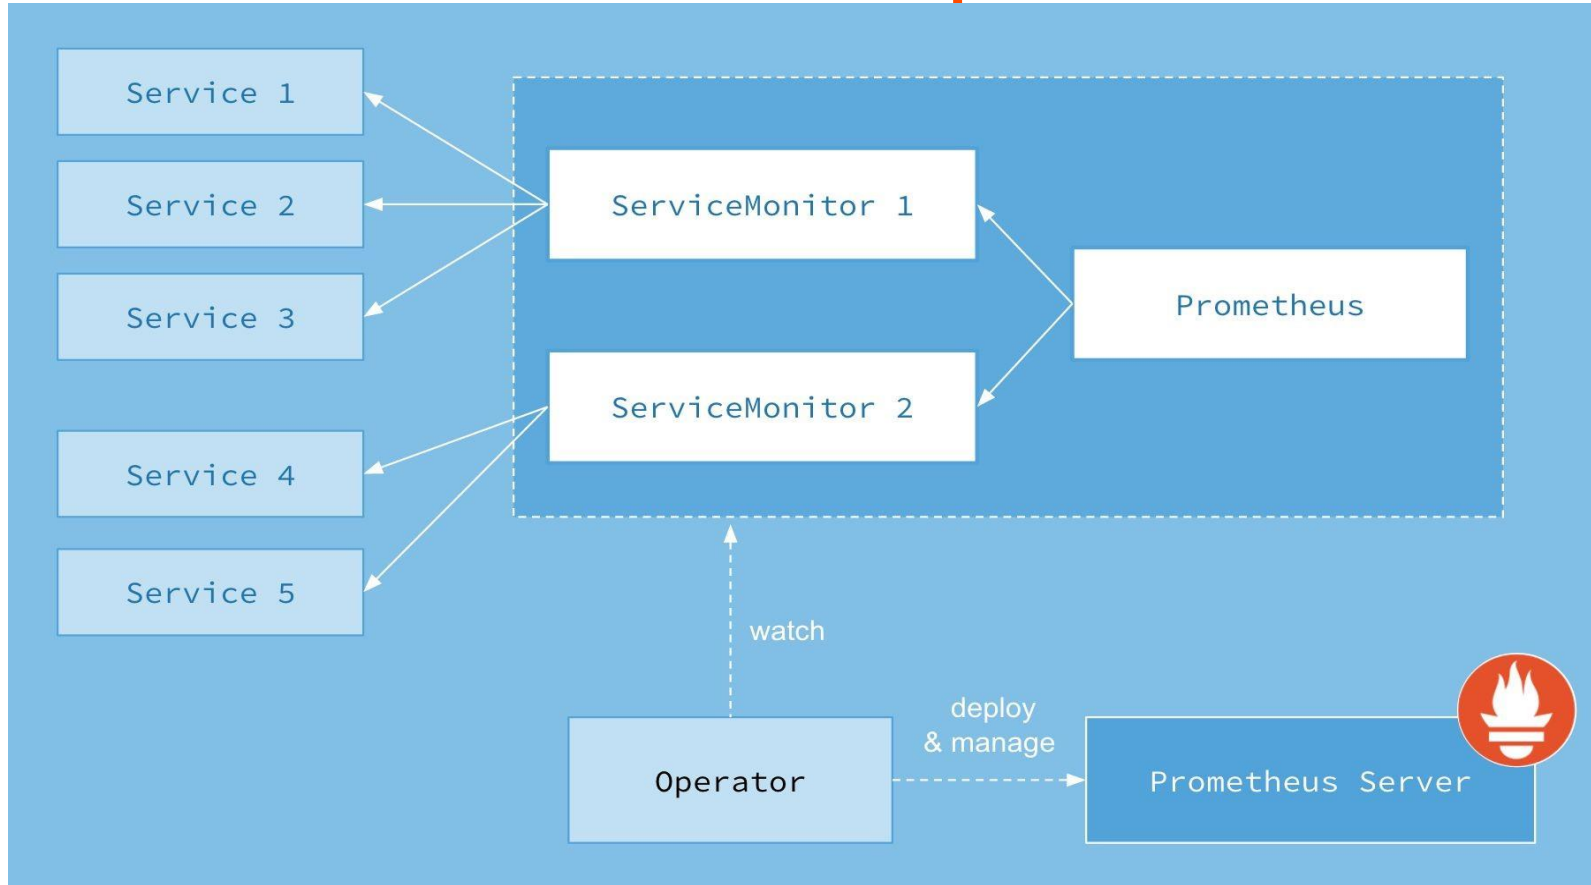
inside of Kubernetes

- What we need
 - Pod specification defining how to run Prometheus
 - Load and manage configuration
 - Service specification to access Prometheus on stable IP
- Options
 - Write own pod+service+petset+... manifests

- Kubernetes Helm chart in the making
<https://github.com/kubernetes/charts/pull/151>
- CoreOS wrote an Operator managing Prometheus and its configuration: <https://github.com/coreos/kube-prometheus>
<https://coreos.com/blog/the-prometheus-operator.html>

scripts/deploy

Prometheus Operator



git checkout 2.install-prometheus

Monitoring Kubernetes infrastructure

scripts/deploy

Configuration and discovery

git checkout 3.monitor-nodes git
checkout 4.monitor-kubernetes
scripts/deploy

git checkout 5.install-grafana

Monitoring services in Kubernetes

scripts/deploy

Configuration and discovery

git checkout 6.monitor-example-app

Practical examples

Queries and dashboards

git checkout 7.add-rules



Prometheus: core-services ▾

System: okidoki ▾

Job: okidoki ▾

Slog ☒

jvmkit (generic) | JVM

jvmkit | MySQL and Memcached

Incoming Request Rate (HTTP & Thrift)

19K rps

99th Percentile Latency

162.6 ms

Error Rate

0.24 rps


Incoming HTTP Request Rate



Incoming HTTP Request Latency



Incoming HTTP Request 5xx Breakdown



Outgoing HTTP Request Rate




Outgoing HTTP Request Latency



Outgoing HTTP Request 5xx Breakdown




Client Failure Accruals



Client Failure Percentage




Client Retries / Requeses




Thrift Successes




Thrift Failures



Thrift 99th Percentile Latency



Further reading



Thank you

Isaac